

```
In [1]: import warnings
# Ignore all warnings
warnings.filterwarnings("ignore")
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
```

## RAIN PREDICTION

Neural Network For  
Classification

### **Rain in Australia**

Predict next-day rain in Australia . . . .

Predicting next-day rain using a dataset containing  
10 years of daily weather observations from  
different locations across Australia.

# TABLE OF CONTENTS

1. IMPORTING LIBRARIES
2. LOADING DATA
3. DATA VISUALIZATION AND CLEANINGS
4. DATA PREPROCESSING
5. MODEL BUILDING
6. CONCLUSION
7. END

## LIBRARIES

### IMPORTING LIBRARIES

```
In [2]: import matplotlib.pyplot as plt
import seaborn as sns
import datetime
from sklearn.preprocessing import LabelEncoder
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import seaborn as sns
from keras.layers import Dense, BatchNormalization, Dropout, LSTM
from keras.models import Sequential
from keras.utils import to_categorical
from keras.optimizers import Adam
from tensorflow.keras import regularizers
from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report, accuracy_score, f1
from keras import callbacks

np.random.seed(0)
```

# LOADING DATA

## LOADING DATA

```
In [3]: data = pd.read_csv("weatherAUS.csv")
data.head()
```

```
Out[3]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	H
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	W	...	71.0	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	NNW	...	44.0	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	W	...	38.0	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	SE	...	45.0	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	ENE	...	82.0	

5 rows × 23 columns

### About the data:

The dataset contains about 10 years of daily weather observations from different locations across Australia. Observations were drawn from numerous weather stations.

In this project, I will use this data to predict whether or not it will rain the next day. There are 23 attributes including the target variable "RainTomorrow", indicating whether or not it will rain the next day or not.

```
In [4]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null object
1   Location              145460 non-null object
2   MinTemp               143975 non-null float64
3   MaxTemp               144199 non-null float64
4   Rainfall              142199 non-null float64
5   Evaporation           82670 non-null float64
6   Sunshine              75625 non-null float64
7   WindGustDir           135134 non-null object
8   WindGustSpeed         135197 non-null float64
9   WindDir9am            134894 non-null object
10  WindDir3pm            141232 non-null object
11  WindSpeed9am          143693 non-null float64
12  WindSpeed3pm          142398 non-null float64
13  Humidity9am           142806 non-null float64
14  Humidity3pm           140953 non-null float64
15  Pressure9am           130395 non-null float64
16  Pressure3pm           130432 non-null float64
17  Cloud9am              89572 non-null float64
18  Cloud3pm              86102 non-null float64
19  Temp9am               143693 non-null float64
20  Temp3pm               141851 non-null float64
21  RainToday             142199 non-null object
22  RainTomorrow          142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB

```

#### Points to notice:

- There are missing values in the dataset
- Dataset includes numeric and categorical values

## DATA VISUALIZATION AND CLEANING

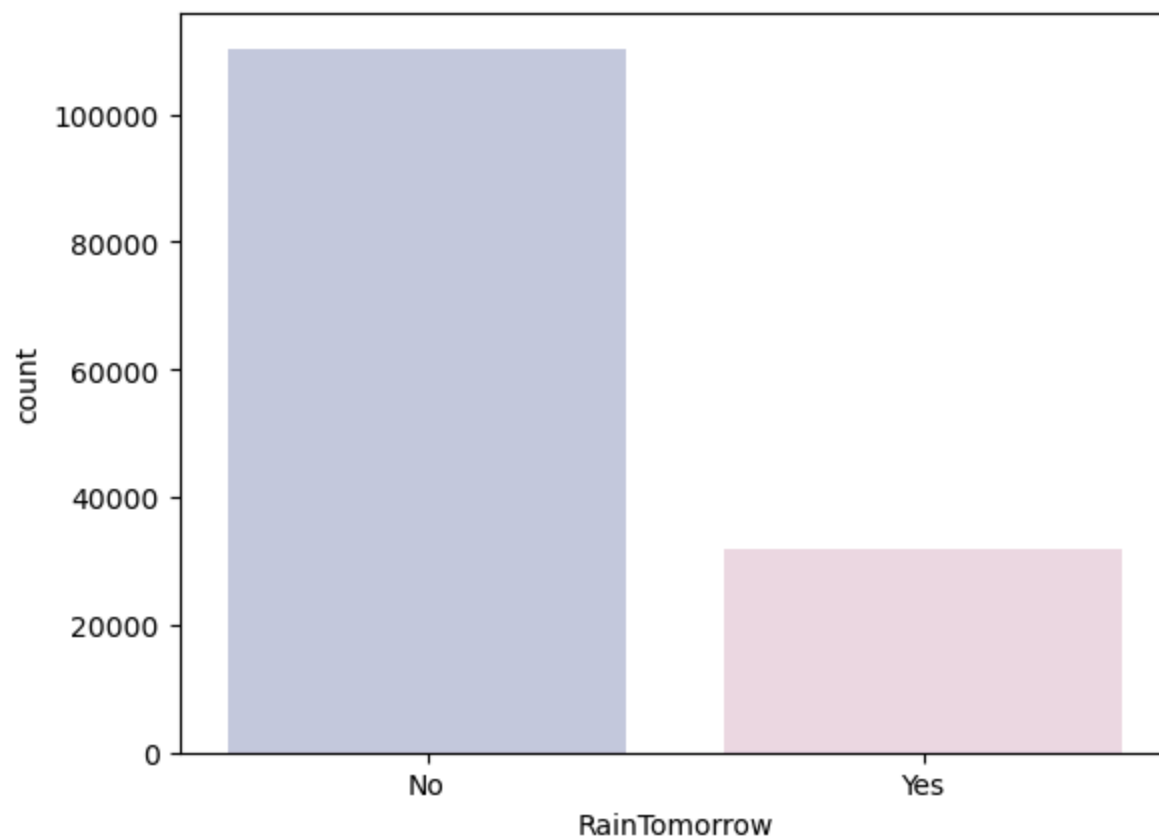
## DATA VISUALIZATION AND CLEANING

**Steps involves in this section:**

- Count plot of target column
- Correlation amongst numeric attributes
- Parse Dates into datetime
- Encoding days and months as continuous cyclic features

```
In [5]: #first of all let us evaluate the target and find out if our data is imbalanced or not  
cols= ["#C2C4E2", "#EED4E5"]  
sns.countplot(x= data["RainTomorrow"], palette= cols)
```

```
Out[5]: <AxesSubplot:xlabel='RainTomorrow', ylabel='count'>
```




```
In [6]: # Correlation amongst numeric attributes  
corrmat = data.corr()  
cmap = sns.diverging_palette(260, -10, s=50, l=75, n=6, as_cmap=True)
```

```
plt.subplots(figsize=(18,18))  
sns.heatmap(corrmat,cmap= cmap,annot=True, square=True)
```

Out[6]: <AxesSubplot:>





### Now I will parse Dates into datetime.

My goal is to build an artificial neural network(ANN). I will encode dates appropriately, i.e. I prefer the months and days in a cyclic continuous feature. As, date and time are inherently cyclical. To let the ANN model know that a feature is cyclical I split it into periodic subsections. Namely, years, months and days. Now for each subsection, I create two new features, deriving a sine transform and cosine transform of the subsection feature.

```
In [7]: #Parsing datetime
#exploring the length of date objects
lengths = data["Date"].str.len()
lengths.value_counts()
```

```
Out[7]: 10    145460
Name: Date, dtype: int64
```

```
In [8]: #There don't seem to be any error in dates so parsing values into datetime
data['Date'] = pd.to_datetime(data["Date"])
#Creating a column of year
data['year'] = data.Date.dt.year

# function to encode datetime into cyclic parameters.
#As I am planning to use this data in a neural network I prefer the months and days in a cyclic continuous feature.

def encode(data, col, max_val):
    data[col + '_sin'] = np.sin(2 * np.pi * data[col]/max_val)
    data[col + '_cos'] = np.cos(2 * np.pi * data[col]/max_val)
    return data

data['month'] = data.Date.dt.month
data = encode(data, 'month', 12)

data['day'] = data.Date.dt.day
data = encode(data, 'day', 31)

data.head()
```



Out[8]:

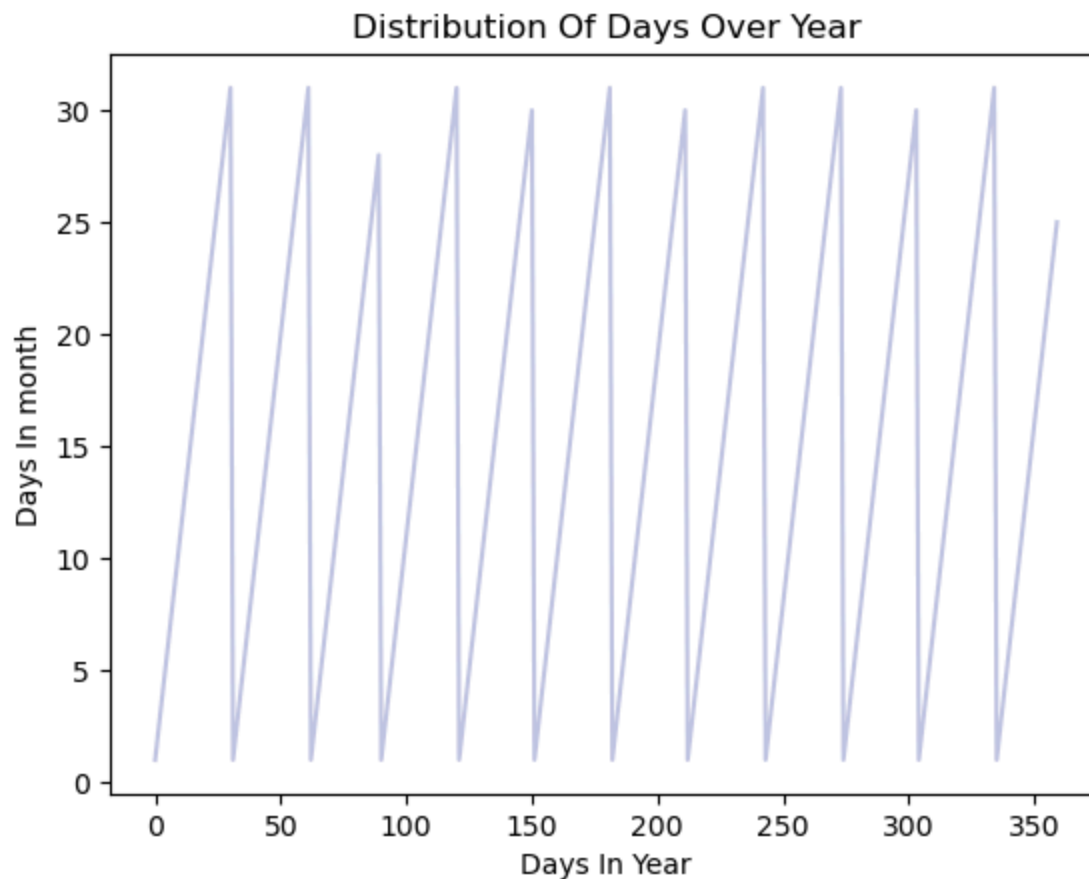
	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Temp3pm	RainT
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	W	...	21.8	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	NNW	...	24.3	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	W	...	23.2	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	SE	...	26.5	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	ENE	...	29.7	

5 rows × 30 columns



In [9]: *# roughly a year's span section*  
 section = data[:360]  
 tm = section["day"].plot(color="#C2C4E2")  
 tm.set\_title("Distribution Of Days Over Year")  
 tm.set\_ylabel("Days In month")  
 tm.set\_xlabel("Days In Year")


Out[9]: Text(0.5, 0, 'Days In Year')



As expected, the "year" attribute of data repeats. However in this for the true cyclic nature is not presented in a continuous manner. Splitting months and days into Sine and cosine combination provides the cyclical continuous feature. This can be used as input features to ANN.


```
In [10]: cyclic_month = sns.scatterplot(x="month_sin",y="month_cos",data=data, color="#C2C4E2")
cyclic_month.set_title("Cyclic Encoding of Month")
cyclic_month.set_ylabel("Cosine Encoded Months")
cyclic_month.set_xlabel("Sine Encoded Months")
```

```
Out[10]: Text(0.5, 0, 'Sine Encoded Months')
```



```
In [11]: cyclic_day = sns.scatterplot(x='day_sin',y='day_cos',data=data, color="#C2C4E2")
cyclic_day.set_title("Cyclic Encoding of Day")
cyclic_day.set_ylabel("Cosine Encoded Day")
cyclic_day.set_xlabel("Sine Encoded Day")
```

```
Out[11]: Text(0.5, 0, 'Sine Encoded Day')
```



Next, I will deal with missing values in categorical and numeric attributes separately

### Categorical variables

- Filling missing values with mode of the column value

```
In [12]: # Get list of categorical variables
s = (data.dtypes == "object")
object_cols = list(s[s].index)

print("Categorical variables:")
print(object_cols)
```

```
Categorical variables:
['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']
```

In [13]: *# Missing values in categorical variables*

```
for i in object_cols:
    print(i, data[i].isnull().sum())
```

Location 0  
WindGustDir 10326  
WindDir9am 10566  
WindDir3pm 4228  
RainToday 3261  
RainTomorrow 3267

In [14]: *# Filling missing values with mode of the column in value*

```
for i in object_cols:
    data[i].fillna(data[i].mode()[0], inplace=True)
```

### Numerical variables

- Filling missing values with median of the column value

In [15]: *# Get list of neumeric variables*

```
t = (data.dtypes == "float64")
num_cols = list(t[t].index)

print("Neumeric variables:")
print(num_cols)
```

Neumeric variables:  
['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am', 'Temp3pm', 'month\_sin', 'month\_cos', 'day\_sin', 'day\_cos']

In [16]: *# Missing values in numeric variables*

```
for i in num_cols:
    print(i, data[i].isnull().sum())
```

```
MinTemp 1485
MaxTemp 1261
Rainfall 3261
Evaporation 62790
Sunshine 69835
WindGustSpeed 10263
WindSpeed9am 1767
WindSpeed3pm 3062
Humidity9am 2654
Humidity3pm 4507
Pressure9am 15065
Pressure3pm 15028
Cloud9am 55888
Cloud3pm 59358
Temp9am 1767
Temp3pm 3609
month_sin 0
month_cos 0
day_sin 0
day_cos 0
```

```
In [17]: # Filling missing values with median of the column in value
```

```
for i in num_cols:
    data[i].fillna(data[i].median(), inplace=True)

data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null  datetime64[ns]
1   Location              145460 non-null  object
2   MinTemp               145460 non-null  float64
3   MaxTemp               145460 non-null  float64
4   Rainfall              145460 non-null  float64
5   Evaporation           145460 non-null  float64
6   Sunshine              145460 non-null  float64
7   WindGustDir           145460 non-null  object
8   WindGustSpeed         145460 non-null  float64
9   WindDir9am            145460 non-null  object
10  WindDir3pm            145460 non-null  object
11  WindSpeed9am          145460 non-null  float64
12  WindSpeed3pm          145460 non-null  float64
13  Humidity9am           145460 non-null  float64
14  Humidity3pm           145460 non-null  float64
15  Pressure9am           145460 non-null  float64
16  Pressure3pm           145460 non-null  float64
17  Cloud9am              145460 non-null  float64
18  Cloud3pm              145460 non-null  float64
19  Temp9am               145460 non-null  float64
20  Temp3pm               145460 non-null  float64
21  RainToday             145460 non-null  object
22  RainTomorrow          145460 non-null  object
23  year                  145460 non-null  int64
24  month                 145460 non-null  int64
25  month_sin             145460 non-null  float64
26  month_cos             145460 non-null  float64
27  day                   145460 non-null  int64
28  day_sin               145460 non-null  float64
29  day_cos               145460 non-null  float64
dtypes: datetime64[ns](1), float64(20), int64(3), object(6)
memory usage: 33.3+ MB

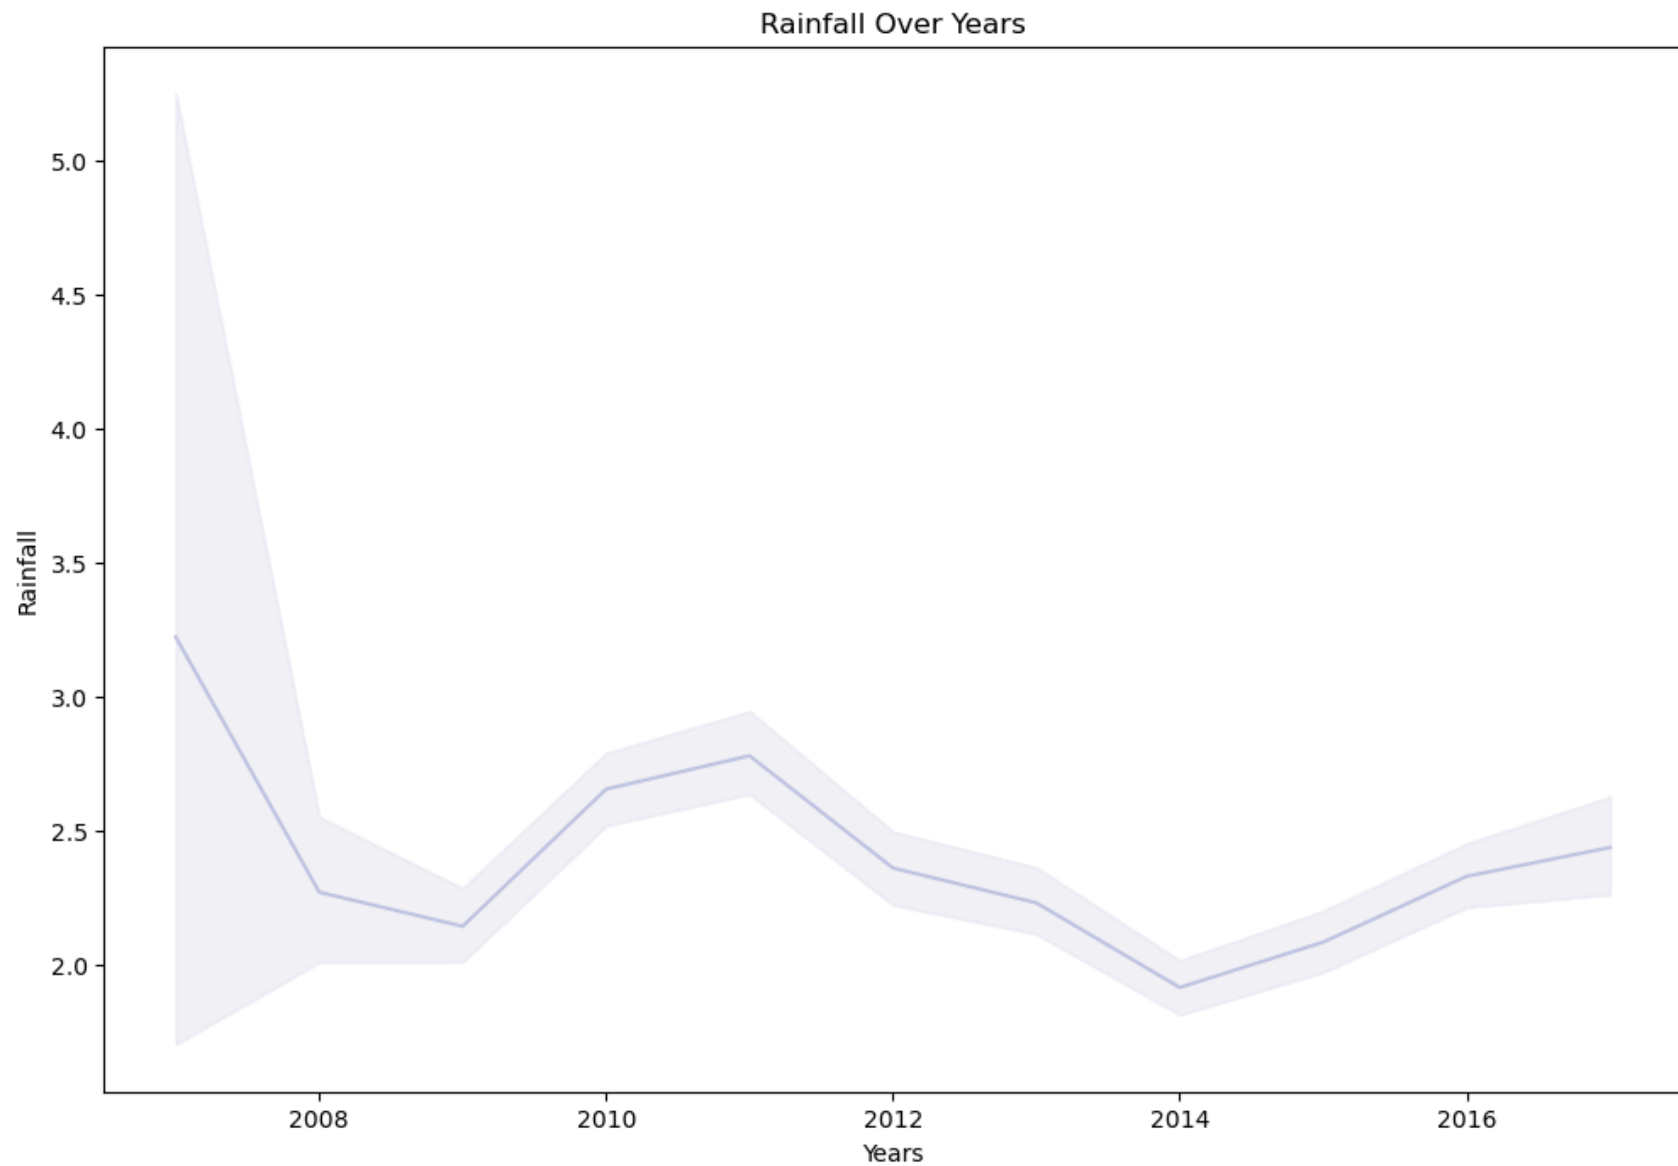
```

```

In [18]: #plotting a lineplot rainfall over years
plt.figure(figsize=(12,8))
Time_series=sns.lineplot(x=data['Date'].dt.year,y="Rainfall",data=data,color="#C2C4E2")
Time_series.set_title("Rainfall Over Years")
Time_series.set_ylabel("Rainfall")
Time_series.set_xlabel("Years")

```

Out[18]: Text(0.5, 0, 'Years')

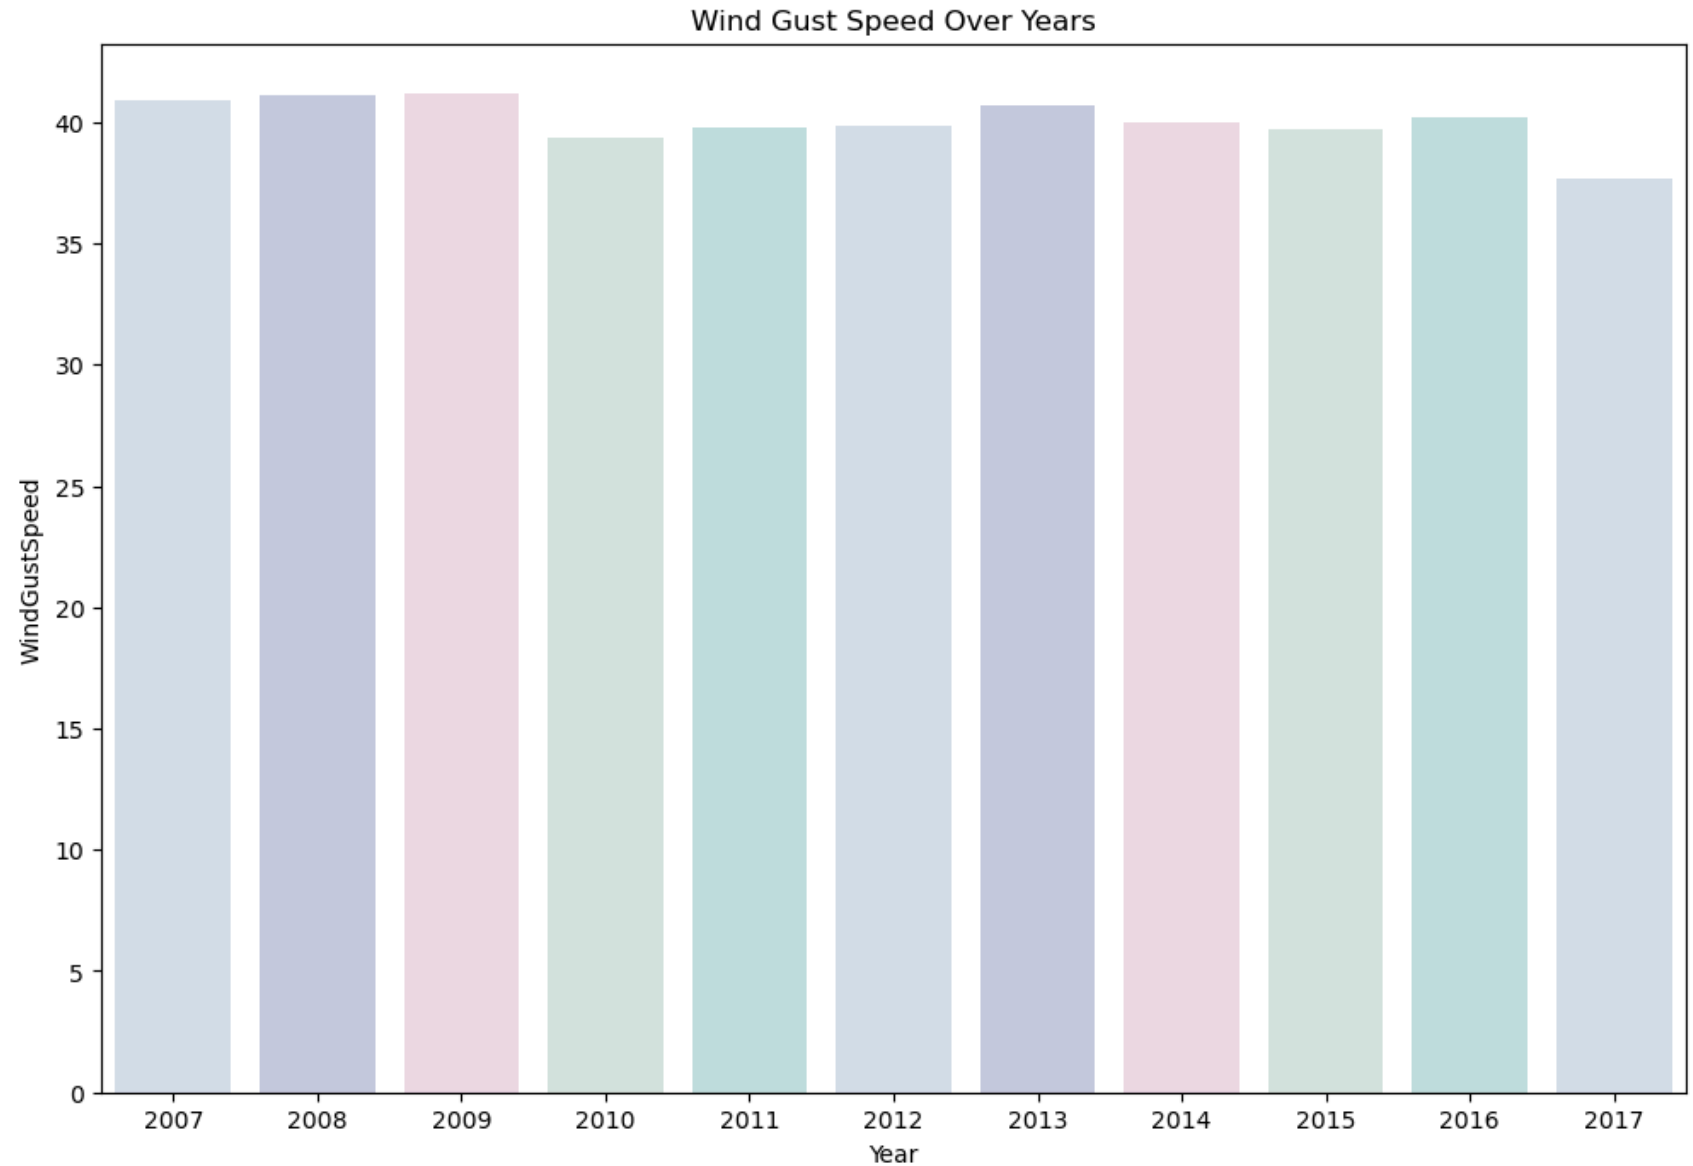


```
In [19]: #Evaluating Wind gust speed over years
colours = ["#D0DBEE", "#C2C4E2", "#EED4E5", "#D1E6DC", "#BDE2E2"]
plt.figure(figsize=(12,8))
Days_of_week=sns.barplot(x=data['Date'].dt.year,y="WindGustSpeed",data=data, ci =None,palette = colours)
Days_of_week.set_title("Wind Gust Speed Over Years")
```



```
Days_of_week.set_ylabel("WindGustSpeed")  
Days_of_week.set_xlabel("Year")
```

Out[19]: Text(0.5, 0, 'Year')



## DATA PREPROCESSING

# DATA PREPROCESSING

## Steps involved in Data Preprocessing:

- Label encoding columns with categorical data
- Perform the scaling of the features
- Detecting outliers
- Dropping the outliers based on data analysis

## Label encoding the catagorical variable

```
In [20]: # Apply Label encoder to each column with categorical data
label_encoder = LabelEncoder()
for i in object_cols:
    data[i] = label_encoder.fit_transform(data[i])

data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   145460 non-null  datetime64[ns]
1   Location                145460 non-null  int32
2   MinTemp                145460 non-null  float64
3   MaxTemp                145460 non-null  float64
4   Rainfall               145460 non-null  float64
5   Evaporation            145460 non-null  float64
6   Sunshine               145460 non-null  float64
7   WindGustDir            145460 non-null  int32
8   WindGustSpeed          145460 non-null  float64
9   WindDir9am             145460 non-null  int32
10  WindDir3pm             145460 non-null  int32
11  WindSpeed9am           145460 non-null  float64
12  WindSpeed3pm           145460 non-null  float64
13  Humidity9am            145460 non-null  float64
14  Humidity3pm            145460 non-null  float64
15  Pressure9am            145460 non-null  float64
16  Pressure3pm            145460 non-null  float64
17  Cloud9am               145460 non-null  float64
18  Cloud3pm               145460 non-null  float64
19  Temp9am                145460 non-null  float64
20  Temp3pm                145460 non-null  float64
21  RainToday              145460 non-null  int32
22  RainTomorrow           145460 non-null  int32
23  year                   145460 non-null  int64
24  month                  145460 non-null  int64
25  month_sin              145460 non-null  float64
26  month_cos              145460 non-null  float64
27  day                    145460 non-null  int64
28  day_sin                145460 non-null  float64
29  day_cos                145460 non-null  float64
dtypes: datetime64[ns](1), float64(20), int32(6), int64(3)
memory usage: 30.0 MB

```

```

In [21]: # Preparing attributes of scale data

features = data.drop(['RainTomorrow', 'Date', 'day', 'month'], axis=1) # dropping target and extra columns

target = data['RainTomorrow']

#Set up a standard scaler for the features

```

```
col_names = list(features.columns)
s_scaler = preprocessing.StandardScaler()
features = s_scaler.fit_transform(features)
features = pd.DataFrame(features, columns=col_names)


features.describe().T
```

Out[21]:

	count	mean	std	min	25%	50%	75%	max
<b>Location</b>	145460.0	-5.633017e-14	1.000003	-1.672228	-0.899139	0.014511	0.857881	1.701250
<b>MinTemp</b>	145460.0	-4.243854e-15	1.000003	-3.250525	-0.705659	-0.030170	0.723865	3.410112
<b>MaxTemp</b>	145460.0	6.513740e-16	1.000003	-3.952405	-0.735852	-0.086898	0.703133	3.510563
<b>Rainfall</b>	145460.0	9.152711e-15	1.000003	-0.275097	-0.275097	-0.275097	-0.203581	43.945571
<b>Evaporation</b>	145460.0	1.352327e-14	1.000003	-1.629472	-0.371139	-0.119472	0.006361	43.985108
<b>Sunshine</b>	145460.0	-4.338304e-15	1.000003	-2.897217	0.076188	0.148710	0.257494	2.360634
<b>WindGustDir</b>	145460.0	1.864381e-14	1.000003	-1.724209	-0.872075	0.193094	1.045228	1.471296
<b>WindGustSpeed</b>	145460.0	-1.167921e-14	1.000003	-2.588407	-0.683048	-0.073333	0.460168	7.243246
<b>WindDir9am</b>	145460.0	-7.433272e-15	1.000003	-1.550000	-0.885669	0.000105	0.885879	1.771653
<b>WindDir3pm</b>	145460.0	1.791486e-15	1.000003	-1.718521	-0.837098	0.044324	0.925747	1.586813
<b>WindSpeed9am</b>	145460.0	-3.422029e-14	1.000003	-1.583291	-0.793380	-0.116314	0.560752	13.086472
<b>WindSpeed3pm</b>	145460.0	1.618238e-14	1.000003	-2.141841	-0.650449	0.037886	0.611499	7.839016
<b>Humidity9am</b>	145460.0	-4.803490e-15	1.000003	-3.654212	-0.631189	0.058273	0.747734	1.649338
<b>Humidity3pm</b>	145460.0	-6.041889e-15	1.000003	-2.518329	-0.710918	0.021816	0.656852	2.366565
<b>Pressure9am</b>	145460.0	2.313398e-14	1.000003	-5.520544	-0.616005	-0.006653	0.617561	3.471111
<b>Pressure3pm</b>	145460.0	4.709575e-15	1.000003	-5.724832	-0.622769	-0.007520	0.622735	3.653960
<b>Cloud9am</b>	145460.0	-2.525820e-14	1.000003	-2.042425	-0.727490	0.149133	0.587445	1.902380
<b>Cloud3pm</b>	145460.0	4.796901e-15	1.000003	-2.235619	-0.336969	0.137693	0.612356	2.036343
<b>Temp9am</b>	145460.0	-3.332880e-15	1.000003	-3.750358	-0.726764	-0.044517	0.699753	3.599302
<b>Temp3pm</b>	145460.0	-2.901899e-15	1.000003	-3.951301	-0.725322	-0.083046	0.661411	3.653834
<b>RainToday</b>	145460.0	1.263303e-14	1.000003	-0.529795	-0.529795	-0.529795	-0.529795	1.887521
<b>year</b>	145460.0	1.663818e-14	1.000003	-2.273637	-0.697391	0.090732	0.878855	1.666978
<b>month_sin</b>	145460.0	1.478014e-15	1.000003	-1.434333	-0.725379	-0.016425	0.692529	1.401483
<b>month_cos</b>	145460.0	4.043483e-16	1.000003	-1.388032	-1.198979	0.023080	0.728636	1.434192
<b>day_sin</b>	145460.0	-6.534944e-18	1.000003	-1.403140	-1.019170	-0.003198	1.012774	1.396744

	count	mean	std	min	25%	50%	75%	max
day_cos	145460.0	-9.540621e-19	1.000003	-1.392587	-1.055520	-0.044639	1.011221	1.455246

```
In [22]: #Detecting outliers
#Looking at the scaled features
colours = ["#D0DBEE", "#C2C4E2", "#EED4E5", "#D1E6DC", "#BDE2E2"]
plt.figure(figsize=(20,10))
sns.boxenplot(data = features,palette = colours)
plt.xticks(rotation=90)
plt.show()
```



```
In [23]: #full data for
features["RainTomorrow"] = target

#Dropping with outlier


features = features[(features["MinTemp"]<2.3)&(features["MinTemp"]>-2.3)]
features = features[(features["MaxTemp"]<2.3)&(features["MaxTemp"]>-2)]
features = features[(features["Rainfall"]<4.5)]
features = features[(features["Evaporation"]<2.8)]
features = features[(features["Sunshine"]<2.1)]
features = features[(features["WindGustSpeed"]<4)&(features["WindGustSpeed"]>-4)]
features = features[(features["WindSpeed9am"]<4)]
features = features[(features["WindSpeed3pm"]<2.5)]
features = features[(features["Humidity9am"]>-3)]
features = features[(features["Humidity3pm"]>-2.2)]
features = features[(features["Pressure9am"]< 2)&(features["Pressure9am"]>-2.7)]
features = features[(features["Pressure3pm"]< 2)&(features["Pressure3pm"]>-2.7)]
features = features[(features["Cloud9am"]<1.8)]
features = features[(features["Cloud3pm"]<2)]
features = features[(features["Temp9am"]<2.3)&(features["Temp9am"]>-2)]
features = features[(features["Temp3pm"]<2.3)&(features["Temp3pm"]>-2)]

features.shape
```

```
Out[23]: (127536, 27)
```

```
In [24]: #Looking at the scaled features without outliers

plt.figure(figsize=(20,10))
sns.boxenplot(data = features,palette = colours)
plt.xticks(rotation=90)
plt.show()
```



Looks Good. Up next is building artificial neural network.

## MODEL BUILDING

# MODEL BUILDING

In this project, we build an artificial neural network.

Following steps are involved in the model building



- Assigning X and y the status of attributes and tags
- Splitting test and training sets
- Initialising the neural network
- Defining by adding layers
- Compiling the neural network
- Train the neural network

```
In [25]: X = features.drop(["RainTomorrow"], axis=1)
y = features["RainTomorrow"]

# Splitting test and training sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)

X.shape
```

```
Out[25]: (127536, 26)
```

```
In [26]: #Early stopping
early_stopping = callbacks.EarlyStopping(
    min_delta=0.001, # minimum amount of change to count as an improvement
    patience=20, # how many epochs to wait before stopping
    restore_best_weights=True,
)


# Initialising the NN
model = Sequential()


# Layers


model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu', input_dim = 26))
model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.25))
model.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))


# Compiling the ANN
opt = Adam(learning_rate=0.00009)
model.compile(optimizer = opt, loss = 'binary_crossentropy', metrics = ['accuracy'])
```


```
# Train the ANN  
history = model.fit(X_train, y_train, batch_size = 32, epochs = 150, callbacks=[early_stopping], validation_split=0.2)
```


Epoch 1/150  
2551/2551  8s 2ms/step - accuracy: 0.7821 - loss: 0.5654 - val\_accuracy: 0.7860 - val\_loss: 0.3927


Epoch 2/150  
2551/2551  6s 2ms/step - accuracy: 0.7842 - loss: 0.4150 - val\_accuracy: 0.7860 - val\_loss: 0.3878


Epoch 3/150  
2551/2551  6s 2ms/step - accuracy: 0.7852 - loss: 0.4112 - val\_accuracy: 0.7860 - val\_loss: 0.3841


Epoch 4/150  
2551/2551  6s 2ms/step - accuracy: 0.7820 - loss: 0.4118 - val\_accuracy: 0.7860 - val\_loss: 0.3821


Epoch 5/150  
2551/2551  6s 2ms/step - accuracy: 0.7847 - loss: 0.4058 - val\_accuracy: 0.7860 - val\_loss: 0.3805


Epoch 6/150  
2551/2551  6s 2ms/step - accuracy: 0.7861 - loss: 0.4017 - val\_accuracy: 0.8433 - val\_loss: 0.3792


Epoch 7/150  
2551/2551  6s 2ms/step - accuracy: 0.8420 - loss: 0.3998 - val\_accuracy: 0.8435 - val\_loss: 0.3775


Epoch 8/150  
2551/2551  6s 2ms/step - accuracy: 0.8405 - loss: 0.3978 - val\_accuracy: 0.8440 - val\_loss: 0.3765


Epoch 9/150  
2551/2551  6s 2ms/step - accuracy: 0.8409 - loss: 0.4007 - val\_accuracy: 0.8430 - val\_loss: 0.3757


Epoch 10/150  
2551/2551  6s 2ms/step - accuracy: 0.8395 - loss: 0.4003 - val\_accuracy: 0.8444 - val\_loss: 0.3748


Epoch 11/150  
2551/2551  6s 2ms/step - accuracy: 0.8418 - loss: 0.3967 - val\_accuracy: 0.8444 - val\_loss: 0.3741


Epoch 12/150  
2551/2551  6s 2ms/step - accuracy: 0.8414 - loss: 0.4011 - val\_accuracy: 0.8449 - val\_loss: 0.3733


Epoch 13/150  
2551/2551  6s 2ms/step - accuracy: 0.8440 - loss: 0.3999 - val\_accuracy: 0.8446 - val\_loss: 0.3729


Epoch 14/150  
2551/2551  5s 2ms/step - accuracy: 0.8427 - loss: 0.3977 - val\_accuracy: 0.8448 - val\_loss: 0.3723


Epoch 15/150  
2551/2551  5s 2ms/step - accuracy: 0.8434 - loss: 0.3988 - val\_accuracy: 0.8450 - val\_loss: 0.3717


Epoch 16/150  
2551/2551  5s 2ms/step - accuracy: 0.8419 - loss: 0.3972 - val\_accuracy: 0.8449 - val\_loss: 0.3710


Epoch 17/150  
2551/2551  5s 2ms/step - accuracy: 0.8438 - loss: 0.3943 - val\_accuracy: 0.8446 - val\_loss: 0.3706

Epoch 18/150  
2551/2551  5s 2ms/step - accuracy: 0.8442 - loss: 0.3950 - val\_accuracy: 0.8447 - val\_loss: 0.3702


Epoch 19/150  
2551/2551  5s 2ms/step - accuracy: 0.8407 - loss: 0.3968 - val\_accuracy: 0.8446 - val\_loss: 0.3700


Epoch 20/150  
2551/2551  5s 2ms/step - accuracy: 0.8447 - loss: 0.3945 - val\_accuracy: 0.8450 - val\_loss: 0.3695


Epoch 21/150  
2551/2551  5s 2ms/step - accuracy: 0.8430 - loss: 0.3942 - val\_accuracy: 0.8453 - val\_loss: 0.3692


Epoch 22/150  
2551/2551  5s 2ms/step - accuracy: 0.8434 - loss: 0.3931 - val\_accuracy: 0.8455 - val\_loss: 0.3690


Epoch 23/150


2551/2551  5s 2ms/step - accuracy: 0.8438 - loss: 0.3924 - val\_accuracy: 0.8449 - val\_loss: 0.3693  
Epoch 24/150


2551/2551  6s 2ms/step - accuracy: 0.8414 - loss: 0.3965 - val\_accuracy: 0.8446 - val\_loss: 0.3692  
Epoch 25/150


2551/2551  5s 2ms/step - accuracy: 0.8417 - loss: 0.3952 - val\_accuracy: 0.8450 - val\_loss: 0.3686  
Epoch 26/150


2551/2551  5s 2ms/step - accuracy: 0.8457 - loss: 0.3892 - val\_accuracy: 0.8442 - val\_loss: 0.3686  
Epoch 27/150


2551/2551  5s 2ms/step - accuracy: 0.8435 - loss: 0.3957 - val\_accuracy: 0.8443 - val\_loss: 0.3685  
Epoch 28/150


2551/2551  5s 2ms/step - accuracy: 0.8447 - loss: 0.3944 - val\_accuracy: 0.8445 - val\_loss: 0.3680  
Epoch 29/150


2551/2551  5s 2ms/step - accuracy: 0.8460 - loss: 0.3889 - val\_accuracy: 0.8441 - val\_loss: 0.3678  
Epoch 30/150


2551/2551  6s 2ms/step - accuracy: 0.8452 - loss: 0.3914 - val\_accuracy: 0.8447 - val\_loss: 0.3681  
Epoch 31/150


2551/2551  6s 2ms/step - accuracy: 0.8435 - loss: 0.3939 - val\_accuracy: 0.8441 - val\_loss: 0.3682  
Epoch 32/150


2551/2551  6s 2ms/step - accuracy: 0.8445 - loss: 0.3951 - val\_accuracy: 0.8446 - val\_loss: 0.3681  
Epoch 33/150


2551/2551  6s 2ms/step - accuracy: 0.8428 - loss: 0.3926 - val\_accuracy: 0.8449 - val\_loss: 0.3680  
Epoch 34/150


2551/2551  6s 2ms/step - accuracy: 0.8397 - loss: 0.3984 - val\_accuracy: 0.8447 - val\_loss: 0.3676  
Epoch 35/150


2551/2551  6s 2ms/step - accuracy: 0.8428 - loss: 0.3949 - val\_accuracy: 0.8447 - val\_loss: 0.3675  
Epoch 36/150


2551/2551  6s 2ms/step - accuracy: 0.8435 - loss: 0.3955 - val\_accuracy: 0.8450 - val\_loss: 0.3676  
Epoch 37/150


2551/2551  5s 2ms/step - accuracy: 0.8445 - loss: 0.3932 - val\_accuracy: 0.8440 - val\_loss: 0.3680  
Epoch 38/150


2551/2551  6s 2ms/step - accuracy: 0.8432 - loss: 0.3927 - val\_accuracy: 0.8445 - val\_loss: 0.3679  
Epoch 39/150


2551/2551  6s 2ms/step - accuracy: 0.8466 - loss: 0.3874 - val\_accuracy: 0.8443 - val\_loss: 0.3674  
Epoch 40/150


2551/2551  6s 2ms/step - accuracy: 0.8438 - loss: 0.3916 - val\_accuracy: 0.8440 - val\_loss: 0.3674  
Epoch 41/150

2551/2551  5s 2ms/step - accuracy: 0.8441 - loss: 0.3900 - val\_accuracy: 0.8446 - val\_loss: 0.3673  
Epoch 42/150

2551/2551  5s 2ms/step - accuracy: 0.8433 - loss: 0.3933 - val\_accuracy: 0.8442 - val\_loss: 0.3674  
Epoch 43/150

2551/2551  5s 2ms/step - accuracy: 0.8432 - loss: 0.3927 - val\_accuracy: 0.8443 - val\_loss: 0.3677  
Epoch 44/150

2551/2551  5s 2ms/step - accuracy: 0.8451 - loss: 0.3937 - val\_accuracy: 0.8444 - val\_loss: 0.3676  
Epoch 45/150

2551/2551  5s 2ms/step - accuracy: 0.8444 - loss: 0.3877 - val\_accuracy: 0.8444 - val\_loss: 0.3674

Epoch 46/150

**2551/2551** ————— 5s 2ms/step - accuracy: 0.8449 - loss: 0.3880 - val\_accuracy: 0.8445 - val\_loss: 0.3674

Epoch 47/150

**2551/2551** ————— 5s 2ms/step - accuracy: 0.8442 - loss: 0.3926 - val\_accuracy: 0.8442 - val\_loss: 0.3671

Epoch 48/150


**2551/2551** ————— 5s 2ms/step - accuracy: 0.8449 - loss: 0.3903 - val\_accuracy: 0.8449 - val\_loss: 0.3670

Plotting training and validation loss over epochs

```
In [27]: history_df = pd.DataFrame(history.history)

plt.plot(history_df.loc[:, ['loss']], "#BDE2E2", label='Training loss')
plt.plot(history_df.loc[:, ['val_loss']], "#C2C4E2", label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc="best")

plt.show()
```




Plotting training and validation accuracy over epochs

```
In [28]: history_df = pd.DataFrame(history.history)

plt.plot(history_df.loc[:, ['accuracy']], "#BDE2E2", label='Training accuracy')
plt.plot(history_df.loc[:, ['val_accuracy']], "#C2C4E2", label='Validation accuracy')

plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



## CONCLUSION

## CONCLUSIONS

### Concluding the model with:

- Testing on the test set
- Evaluating the confusion matrix
- Evaluating the classification report


```
In [29]: # Predicting the test set results
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)
```

798/798 ————— 1s 1ms/step

```
In [30]: # confusion matrix
cmap1 = sns.diverging_palette(260,-10,s=50, l=75, n=5, as_cmap=True)
plt.subplots(figsize=(12,8))
cf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(cf_matrix/np.sum(cf_matrix), cmap = cmap1, annot = True, annot_kws = {'size':15})
```

Out[30]: <AxesSubplot:>





```
In [31]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.94	0.91	20110
1	0.69	0.51	0.59	5398
accuracy			0.85	25508
macro avg	0.78	0.72	0.75	25508
weighted avg	0.84	0.85	0.84	25508

In [ ]:

In [ ]: